

Capability Scorecard

CareerForge Day 7 — Engineering
Defence Board

Student: Vibhuti Uttarde | 3rd Year Diploma | Information Technology | Vidyalkar Polytechnic

Self-score before the live defense. Founder/evaluator scoring may differ.

Capability Summary Table

Capability	Evidence From The Week	Self-Rating	Defense Risk
Debugging Discipline	Day 1: Found 4 stacked bugs (port, env URL, JWT expiry, Content-Type header) using netstat, Network tab, and source-code reading. Rejected founder framing. Fixed in one commit. Full workflow verified before/after.	Strong	Low — commit IDs, log evidence, and before/after workflow screenshots all present.
Systems Reasoning	Day 2: PostgreSQL write latency identified as origin failure at 10:01:10, before all 5 downstream symptoms. 4 hypotheses evaluated with log-line evidence. Causal chain walked from DB to every reported symptom.	Strong	Low — log sequence is unambiguous and can be walked step by step from memory.
Production Verification	Day 3: 3 cascading .env blockers found on same code path. Declared critical workflow first. Demo blocked with specific evidence. All 4 fixes were .env only — no source code changed.	Adequate	Medium — only c_101 tested, not all 4 candidates. Gap was acknowledged in the report.
Ambiguity Ownership	Day 4: 5 ambiguities surfaced in one founder sentence before any code written. One clarification question unlocked 4 answers. v1 scoped to shortlist button + eligibility warning + server log only.	Adequate	Medium — evaluator may probe whether eligibility compliance risk was fully resolved, or whether one question was sufficient.
AI Engineering Judgment	Day 5: 3 of 6 AI recommendations blocked with specific failure sequences. 1 modified with atomic constraint fix. 2 trusted. Full AI module rejected on Day 4. Correct outcomes on all AI tasks.	Adequate	Medium — correct outcomes reached, but failure-sequence evaluation was sometimes a recovery move rather than a first move.
Pressure Prioritization	Day 6: Checkout paused T+01 before full harm quantified. Engineering directed to root cause (webhook idempotency), not dashboard. AI rollback recommendation modified. Narrow refund scope. Proactive comms before public pressure peaked.	Strong	Low — priority order held under founder pressure with stated rationale. One real mistake (banner vs FAQ) acknowledged proactively.
Engineering Defense	Day 7: Full dossier with per-day evidence bullets, honest capability self-assessment, growth prescription with 4-week plan, and self-identified gaps. No defensive language used.	Adequate	Medium — written defense is strong. Live viva tests whether reasoning holds under direct questioning without the report.

Strongest Demonstrated Capability

Capability: Debugging Discipline and Systems Reasoning (tied).

Evidence: Both Days 1 and 2 show the identical core behaviour — reject surface framing, follow the evidence chain, act on the origin not the symptom. Day 1 in source code. Day 2 in system logs. The method is identical and transferable to any domain.

Most Exposed Capability

Capability: AI Engineering Judgment.

Evidence: I defaulted to trusting AI throughout this challenge and developed the correct evaluation habit by working through the failure sequences myself. The Day 5 Redis recommendation required manual crash-scenario reconstruction to find the flaw. That should be a first-move check, not a recovery move.

Defense Preparation Notes

Be ready to defend live:

- Day 3: Why I blocked with only c_101 tested — and what the remaining risk was.
- Day 4: Why one clarification question was sufficient — and what I would do if the founder's answer turned out to be wrong post-pilot.
- Day 5: Walk the Redis failure sequence from memory — stock=1, A and C both hold paid orders, Postgres never updated.
- Day 6: Defend the bounded window decision at T+06 — why I did not fully stop the drop even though founder pressure was rising.
- All days: Why I consistently challenged founder and AI framing rather than accepting it.